
LiveTS: A Leakage-Proof, Continuously-Refreshed Benchmark for Zero-Shot Time Series Forecasting

Anonymous Authors

(LiveTS working draft; author list to be finalized)

Abstract

Time series foundation models (TSFMs) are now routinely evaluated in the “zero-shot” regime on benchmarks that predate their pretraining corpora. This creates two failure modes that no amount of statistical care can repair post hoc: (i) *pretraining contamination* — the test data may literally be in the training set — and (ii) *community overfitting* — static test sets are reused across thousands of experiments, turning reported gains into selection effects. We introduce **LiveTS**, a benchmark that makes both failure modes physically impossible rather than statistically unlikely. LiveTS continuously collects time series from six public real-time domains (energy, weather, air quality, traffic, crypto/FX, and web traffic) with point-in-time snapshot semantics, registers a conservative public release date for every model, and scores a model only on forecast windows whose evaluation cutoff falls *after* that release date (future-only / as-of evaluation). We pre-registered the full protocol — leakage boundaries, metrics (MASE/CRPS/WQL), geometric-mean aggregation with bootstrap confidence intervals, and Diebold–Mariano tests with Holm correction — before running the main experiments. Using historical simulation over 205 daily series and three rolling cutoffs (50,284 forecast windows), we report the first *clean* scores for six open TSFMs against statistical baselines and supervised baselines retrained strictly before each cutoff. We find no universal winner: TSFMs deliver consistent gains on seasonal domains but fail to beat a seasonal-naïve baseline on financial series, and rankings shift across cutoffs — evidence that single-number leaderboards on static benchmarks overstate what zero-shot forecasting has achieved. LiveTS ships with an open collection pipeline, audit-ready snapshot hashes, and a rolling leaderboard that refreshes monthly, so that the benchmark cannot be overfit faster than reality produces new ground truth.

1 Introduction

Benchmark evaluation in time series forecasting is facing a quiet crisis. The field’s standard test sets — ETT, Electricity, Traffic, Weather, the Monash archive [Godahewa et al., 2021] — were assembled years before the current generation of time series foundation models (TSFMs) was trained [Ansari et al., 2024, Das et al., 2024, Woo et al., 2024, Shi et al., 2025]. Modern TSFMs are pretrained on web-scale corpora of public time series; the very datasets used to certify their “zero-shot” ability are, at least in part, plausibly inside their training distribution, and in the worst case inside their training set verbatim [Meyer et al., 2025]. Reported zero-shot numbers are therefore upper bounds of unknown looseness.

A second, older problem compounds the first. Static test sets are evaluated millions of times by a community that collectively tunes architectures, hyperparameters, and random seeds against them. Even with honest individual practice, the community-level process is adaptive data reuse, and the resulting benchmark gains need not transfer.

Both problems share a root cause: *the test data existed before the models did*. We argue the fix must be structural, not statistical. **LiveTS** evaluates each model only on data generated after the model’s weights were publicly released. Contamination and test-set overfitting are then impossible by construction — a model cannot have been pretrained on data that did not exist, and the community cannot have tuned against ground truth that had not yet been produced.

Our contributions:

1. **A leakage-proof evaluation design** (§3): future-only / as-of evaluation with registered model release dates, rolling cutoffs, point-in-time (PIT) collection semantics, and hashed raw snapshots for audit.
2. **A pre-registered protocol** (§3.3, full text in Appendix A): metrics, aggregation, uncertainty quantification, significance testing, and an amendment policy, frozen before the main experiments.
3. **First clean scores** (§4): a historical-simulation study over 205 daily series from six domains, three rolling cutoffs, six open TSFMs, statistical baselines (seasonal naive, AutoETS), and supervised baselines (DLinear, PatchTST, iTransformer) retrained strictly before each cutoff — 50,284 forecast windows with full JSONL provenance.
4. **A living benchmark** (§5): an open-source collection pipeline running daily across redundant free sources, and a monthly-refresh leaderboard whose ground truth is produced by reality after submissions close.

2 Related Work

Static TSFM benchmarks. GIFT-Eval [Aksu et al., 2024] and fev-bench [Shchur et al., 2025] aggregate large collections of public datasets and standardize protocol details, but both are static: their test sets predate the models under evaluation, so pretraining overlap must be handled by best-effort de-duplication, which cannot rule out paraphrased or derived copies [Meyer et al., 2025]. The Monash archive [Godahewa et al., 2021] is the canonical pre-TSFM benchmark and is now saturated. LiveTS is complementary: it trades corpus breadth for a guarantee no static benchmark can offer.

Live evaluation in other fields. The M6 competition [Makridakis et al., 2025] ran a genuine live forecasting evaluation but was a one-off event in a single (financial) domain. LiveBench [White et al., 2025] and dynamic LLM leaderboards apply the same future-only insight to language models. LiveTS brings this paradigm to multi-domain time series with a reproducible open pipeline.

Evaluation methodology. A line of work documents protocol inconsistencies (look-back lengths, normalization, per-dataset tuning) that make cross-paper comparisons unreliable [Meyer et al., 2025]. LiveTS fixes a single pre-registered protocol and forbids per-dataset tuning; its methodological contribution is the combination of PIT data semantics with registered release dates.

Contamination auditing. Membership-inference and overlap-detection methods estimate contamination post hoc. LiveTS gives the complementary counterfactual: the gap between a model’s score on old benchmarks and its clean LiveTS score is an upper bound on what contamination plus community overfitting bought (we develop this analysis jointly with a companion contamination-audit project).

3 LiveTS Design

3.1 Data layer

Six domains, chosen for public availability, daily-or-finer update frequency, key-free access, and at least two redundant sources each: electric grid load/price (NESO UK, SMARD DE), weather station aggregates (Open-Meteo), PM2.5 air quality (Open-Meteo AQ), public transit ridership (NYC MTA/Socrata), crypto and FX rates (Coinbase, ECB/Frankfurter), and Wikipedia pageviews (Wikimedia). A configurable collector runs daily via cron on two geographically separate hosts, stores the raw API response verbatim, tags every record with `collected_at` (PIT semantics), and archives snapshots to object storage with a SHA-256 manifest so that any later revision of source

data is detectable. The evaluation set currently spans **205 daily series** (weather 80, web traffic 58, crypto/FX 34, air quality 20, traffic 7, energy 6); the series list is configuration, not code, and grows without changing the protocol.

3.2 Leakage boundary: release dates and future-only scoring

Every evaluated model registers a **release date**: the earliest publicly verifiable timestamp of its weights (we use the Hugging Face repository creation date; when in doubt, the earlier date is used — conservative in the direction of *fewer* clean windows). Given rolling cutoffs $c_1 < c_2 < \dots$, a model’s clean score aggregates only forecast windows with cutoff $>$ release date. Inputs at each forecast origin are strictly pre-origin observations; targets are the subsequent horizon. In live rounds the target values physically do not exist at submission time; the historical-simulation study in §4 replays the same rule against archived history, which is exact provided the model’s weights were frozen at its release date.

3.3 Pre-registered protocol

Frozen 2026-08-09 (summary; full text Appendix A). Three cutoffs (2025-01-01, 2025-07-01, 2026-01-01); four forecast origins per cutoff drawn from the 180 days following it; horizon 14 at daily frequency; look-back = each model’s maximum context up to 2048 points, identical across domains; no per-dataset tuning of any kind. Metrics: MASE [Hyndman and Koehler, 2006] (seasonal in-sample naive scale; period 7 for seasonal domains, 1 for financial), CRPS approximated from nine deciles [Gneiting and Raftery, 2007], and weighted quantile loss (WQL). Aggregation: arithmetic mean over origins $\times$ seeds within a series, then geometric mean across series; 95% CIs from a cross-series bootstrap ($B=1000$). Sampling-based models run seeds $\{0, 1, 2\}$; deterministic models are flagged. Significance: paired bootstrap and Diebold–Mariano [Diebold and Mariano, 1995] with the Harvey–Leybourne–Newbold small-sample correction [Harvey et al., 1997] on shared windows, Holm-corrected [Holm, 1979] at $\alpha=0.05$. Point-only models receive no probabilistic score (we do not impute quantiles from point forecasts). Supervised baselines are retrained from scratch on strictly pre-cutoff data for every cutoff, making every cutoff clean for them by construction.

3.4 Service layer

Snapshots and manifests live in object storage (Cloudflare R2); a monthly evaluation round opens after each cutoff, accepts one submission per model per round, and publishes scores when the target window has fully materialized. New models must register a release date before their first round; protocol changes are timestamped amendments that never retroactively alter published scores.

4 Historical-Simulation Experiments

4.1 Setup

205 series $\times$ 3 cutoffs $\times$ 4 origins $\times$ horizon 14 ($\approx 2,460$ windows per deterministic model; $\times 3$ seeds for sampling models and supervised baselines). Models: seasonal naive; AutoETS; DLinear [Zeng et al., 2023], PatchTST [Nie et al., 2023], iTransformer [Liu et al., 2024] (global models, retrained pre-cutoff, multi-quantile loss, 3 seeds); Chronos-Bolt-small/base and Chronos-T5-small [Ansari et al., 2024]; Time-MoE-50M [Shi et al., 2025] (point-only); Moirai-1.1-R-small [Woo et al., 2024]; TimesFM-2.5-200M [Das et al., 2024]. All zero-shot inference used a single configuration per model across all domains. Every result row carries the git commit, environment versions, and run timestamp (JSONL, open).

4.2 Main results

Table 1 and Figure 1 report the main results; Table 2 and Figure 2 break them down by domain. Three observations:

1. **Release-date gating matters.** TimesFM-2.5’s clean score covers only the cutoff after its 2025-09 release; comparing it to models with three clean cutoffs on the pooled table would

Table 1: Clean scores on LiveTS (historical simulation, 205 series, future-only cutoffs). Geometric mean across series of per-series mean MASE/WQL; 95% cross-series bootstrap CI ($B=1000$). [†] fewer than three clean cutoffs due to release-date gating — compare on overlapping cutoffs only. Supervised baselines (release “per-cutoff”) are retrained strictly pre-cutoff, so every cutoff is clean for them.

Model	Release	geo-MASE [95% CI]	geo-WQL	Series	Windows
timesfm-2.5-200m [†]	2025-09-02	0.789 [0.711, 0.873]	0.120	205	820
chronos-bolt-small	2024-11-25	0.893 [0.817, 0.976]	0.127	207	2484
chronos-bolt-base	2024-11-25	0.895 [0.826, 0.980]	0.127	207	2484
chronos-t5-small	2024-02-21	0.956 [0.877, 1.048]	0.132	205	7380
moirai-1.1-r-small	2024-06-14	0.965 [0.880, 1.051]	0.140	205	7380
time-moe-50m	2024-09-21	1.004 [0.923, 1.111]	—	205	2460
patchtst	per-cutoff	1.099 [1.019, 1.185]	0.179	207	7452
autoets	—	1.133 [1.050, 1.223]	0.212	205	2460
seasonal_naive	—	1.181 [1.092, 1.268]	0.200	205	2460
dlinear	per-cutoff	1.597 [1.428, 1.788]	0.334	207	7452
itransformer	per-cutoff	4.220 [3.860, 4.627]	0.933	207	7452

Table 2: Per-domain geo-MASE. No model wins everywhere; on crypto/FX no model beats seasonal naive except TimesFM-2.5 on its single clean cutoff.

Model	air_quality	crypto_fx	energy	traffic	weather	web_traffic
timesfm-2.5-200m	0.75	2.16	0.93	0.75	0.68	0.54
chronos-bolt-small	0.79	2.66	0.90	1.24	0.70	0.67
chronos-bolt-base	0.79	2.75	0.89	1.19	0.70	0.67
chronos-t5-small	0.83	2.98	0.97	1.46	0.75	0.68
moirai-1.1-r-small	0.83	2.75	1.03	1.61	0.76	0.72
time-moe-50m	0.81	3.60	0.89	1.04	0.80	0.71
patchtst	0.85	2.98	1.02	1.28	0.95	0.82
autoets	1.17	2.57	1.25	1.39	0.93	0.88
seasonal_naive	1.04	2.57	1.30	1.71	1.05	0.87
dlinear	1.02	8.01	1.17	1.46	1.25	1.08
itransformer	2.74	13.61	3.16	9.00	3.46	3.08

be flattering nonsense. LiveTS makes this visible instead of averaging it away. On shared windows, TimesFM-2.5 is significantly better than every other model after Holm correction (Appendix C).

2. **No universal winner.** On financial series (crypto/FX), no TSFM beats seasonal naive on MASE — consistent with near-martingale dynamics — while on seasonal domains (weather, air quality, traffic, web) TSFMs deliver 10–40% MASE improvements (Figure 2).
3. **Significance survives correction only for the largest gaps.** After Holm correction, Chronos-family and Moirai models are significantly better than seasonal naive; PatchTST is statistically indistinguishable from seasonal naive and AutoETS on shared windows; several pairwise TSFM comparisons that look decisive as point estimates are not significant (Appendix C).

Supervised baselines were retrained pre-cutoff with a fixed generic configuration (multi-quantile loss, input size 96, 1000 steps, no per-dataset tuning). PatchTST is competitive with weaker TSFMs; DLinear and iTransformer underperform under this no-tuning regime. iTransformer runs in univariate mode ($n_{\text{series}}=1$), a conservative configuration at odds with its multivariate design; we report it as a lower bound and note this as a limitation rather than tuning it post hoc.

4.3 Robustness

Figure 3 shows per-cutoff scores: mid-table rankings (Chronos-T5 vs. Moirai vs. Time-MoE vs. PatchTST) reorder across cutoffs, while the top (Chronos-Bolt) and bottom are stable. Seed variance for sampling models is small relative to cross-series dispersion. This is quantitative evidence that a single static split materially misranks models — and an argument for rolling evaluation as the default.




Figure 1: Clean geo-MASE with 95% bootstrap CIs. Grey bar: model with fewer clean cutoffs (release-date gating); dashed line: seasonal-naive parity.




Figure 2: Per-domain geo-MASE relative to seasonal naive ($\log_2$ ratio; blue = better than naive). TSFMs win on seasonal domains and lose on crypto/FX.

4.4 Clean vs. legacy scores (contamination gain)

For each TSFM we will contrast its published results on legacy benchmarks with its LiveTS clean score relative to identical baselines; this analysis is being developed jointly with a companion contamination-audit project and is deferred to the camera-ready.

5 The Live Benchmark

Round-0 opens with the protocol frozen in Appendix A. Monthly cadence: cutoff on the 1st, submissions within 7 days (forecasts only — no post-cutoff data access), scoring after the 14-day target window completes, publication with full JSONL provenance. Anti-gaming: one submission per




Figure 3: Per-cutoff geo-MASE. Mid-table rankings shift across cutoffs; a single static split would misrank models.

model per round; release-date registration precedes first participation; raw snapshot hashes published for third-party audit; the aggregation code is deterministic and open.

6 Limitations and Ethics

Univariate daily-frequency focus (intraday and multivariate tracks are future work); free-tier sources can rate-limit or discontinue (mitigated by ≥ 2 sources per domain, but a domain could still degrade); public infrastructure data skews toward well-instrumented regions; release dates rely on public timestamps and may be later than private availability (our conservative rule errs toward fewer clean windows, not more). Supervised baselines use one generic configuration — their scores are lower bounds, not the best achievable. The benchmark evaluates forecast quality only; downstream decision costs are out of scope.

References

- Taha Aksu, Gerald Woo, Juncheng Liu, Xu Liu, Chenghao Liu, Silvio Savarese, Caiming Xiong, and Doyen Sahoo. GIFT-Eval: A benchmark for general time series forecasting model evaluation. *arXiv preprint arXiv:2410.10393*, 2024.
- Abdul Fatir Ansari, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Sundar Rangapuram, Sebastian Pineda Arango, Shubham Kapoor, Jasper Zschiegner, Danielle C. Maddix, Michael W. Mahoney, Kari Torkkola, Andrew Gordon Wilson, Michael Bohlke-Schneider, and Yuyang Wang. Chronos: Learning the language of time series. *Transactions on Machine Learning Research*, 2024. arXiv:2403.07815.
- Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. A decoder-only foundation model for time-series forecasting. In *International Conference on Machine Learning (ICML)*, 2024. arXiv:2310.10688.
- Francis X. Diebold and Roberto S. Mariano. Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3):253–263, 1995.
- Tilmann Gneiting and Adrian E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007.
- Rakshitha Godahewa, Christoph Bergmeir, Geoffrey I. Webb, Rob J. Hyndman, and Pablo Montero-Manso. Monash time series forecasting archive. In *Neural Information Processing Systems Track on Datasets and Benchmarks*, 2021.

- David Harvey, Stephen Leybourne, and Paul Newbold. Testing the equality of prediction mean squared errors. *International Journal of Forecasting*, 13(2):281–291, 1997.
- Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.
- Rob J. Hyndman and Anne B. Koehler. Another look at measures of forecast accuracy. *International Journal of Forecasting*, 22(4):679–688, 2006.
- Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, and Mingsheng Long. iTransformer: Inverted transformers are effective for time series forecasting. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.06625.
- Spyros Makridakis, Evangelos Spiliotis, Ross Hollyman, Fotios Petropoulos, Norman Swanson, and Anil Gaba. The M6 forecasting competition: Bridging the gap between forecasting and investment decisions. *International Journal of Forecasting*, 41(4):1315–1354, 2025.
- Marcel Meyer, Sascha Kaltenpoth, Kevin Zalipski, and Oliver Müller. Rethinking evaluation in the era of time series foundation models: (un)known information leakage challenges. *arXiv preprint arXiv:2510.13654*, 2025.
- Yuqi Nie, Nam H. Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2211.14730.
- Oleksandr Shchur, Abdul Fatir Ansari, Caner Turkmen, Lorenzo Stella, Nick Erickson, Pablo Guerron, Michael Bohlke-Schneider, and Yuyang Wang. fev-bench: A realistic benchmark for time series forecasting. *arXiv preprint arXiv:2509.26468*, 2025.
- Xiaoming Shi, Shiyu Wang, Yuqi Nie, Dianqi Li, Zhou Ye, Qingsong Wen, and Ming Jin. Time-MoE: Billion-scale time series foundation models with mixture of experts. In *International Conference on Learning Representations (ICLR)*, 2025. arXiv:2409.16040.
- Colin White, Samuel Dooley, Manley Roberts, Arka Pal, Benjamin Feuer, Siddhartha Jain, Ravid Shwartz-Ziv, Neel Jain, Khalid Saifullah, Sreemanti Dey, Shubh-Agrawal, Sandeep Singh Sandha, Siddhartha Naidu, Chinmay Hegde, Yann LeCun, Tom Goldstein, Willie Neiswanger, and Micah Goldblum. LiveBench: A challenging, contamination-limited LLM benchmark. In *International Conference on Learning Representations (ICLR)*, 2025.
- Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Unified training of universal time series forecasting transformers. In *International Conference on Machine Learning (ICML)*, 2024. arXiv:2402.02592.
- Ailing Zeng, Muxi Chen, Lei Zhang, and Qiang Xu. Are transformers effective for time series forecasting? In *AAAI Conference on Artificial Intelligence*, 2023. arXiv:2205.13504.

A Pre-registered protocol v1.0 + amendment A1

The full pre-registration (frozen 2026-08-09) is distributed with the benchmark repository (docs/protocol-prereg.md); amendment A1 (series expansion 25→205 and supervised baselines) was registered before the expanded runs and does not alter previously published results.

B Data sources and series list

See docs/data-sources.md and the generated series inventory in the repository. Every series has ≥ 800 daily points (≥ 700 for FX, trading-day calendar).

C Full significance matrices

The full pairwise matrix (56 pairs, paired bootstrap $B=10,000$ and DM–HLN, Holm-corrected) is generated by scripts/significance.py and distributed as docs/significance.md.

D Reproduction

One command per table; JSONL schema; environment lockfiles; commit hashes. See `REPRODUCIBILITY.md` in the repository.